

cal_housing

November 16, 2020

0.0.1 Linear Regression Modeling of California Home Prices

0.0.2 University of California, Santa Barbara

0.0.3 PSTAT 135/235: Big Data Analytics

0.0.4 Last Updated: May 30, 2020

TOTAL POINTS: 10

Instructions

In this project, you will work with the California Home Price dataset to train a regression model and predict median home prices. Please do the following:

- 1) (6 PTS) Go through all code and fill in the missing cells. This will prep data, train a model, predict, and evaluate model fit. Compute and report the Mean Squared Error (MSE).
- 2) (1 PT) Repeat Part 1 with at least one additional feature from the original set.
- 3) (2 PTS) Repeat Part 1 with at least one engineered feature based on one or more variables from the original set.
- 4) (1 PT) Repeat Part 1 using Lasso Regression

Please report results in the following way:

In the **RESULTS SECTION** table at the very bottom, there are three cells where you should copy your code from parts 2,3,4.

In the very last cell, print a dataframe containing two columns: `question_part` and `MSE`.

This dataframe must report your MSE results.

Data Source

StatLib—Datasets Archive

<http://lib.stat.cmu.edu/datasets/>

```
[1]: import os
import pandas as pd

from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()
```

```
[2]: # read text file into pyspark dataframe
filename = 'cal_housing_data_preproc_w_header.txt'
df = spark.read.csv(filename, inferSchema=True, header = True)
```

```
[3]: df.show(3)
```

```
+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+
|median_house_value|    median_income|housing_median_age|total_rooms|total_bedro
oms|population|households|latitude|longitude|
+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+
|          452600.0|          8.3252|          41.0|          880.0|
129.0|          322.0|          126.0|          37.88|         -122.23|
|          358500.0|          8.3014|          21.0|          7099.0|
1106.0|          2401.0|          1138.0|          37.86|         -122.22|
|          352100.0|7.257399999999999|          52.0|          1467.0|
190.0|          496.0|          177.0|          37.85|         -122.24|
+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+
only showing top 3 rows
```

0.0.5 Additional Preprocessing

We want to do three more things before training a model:

SCALING (1 POINT)

Scale the response variable median_house_value, dividing by 100000 and saving into column median_house_value_final

```
[4]: df_scaled = df.withColumn("median_house_value_final", df.median_house_value / 100000)
```

FEATURE ENGINEERING (1 POINT)

Add new feature: rooms_per_household

```
[5]: df2 = df_scaled.withColumn("rooms_per_household", df.total_rooms / df.households)
```

SELECT AND STANDARDIZE FEATURES (2 POINTS)

```
[6]: # retain these predictors for Part 1
vars_to_keep = ["median_house_value_final",
                "total_bedrooms",
                "population",
                "households",
                "median_income",
```

```

        "rooms_per_household"]

from pyspark.ml.feature import VectorAssembler
# subset the dataframe on these predictors
assembler = VectorAssembler(inputCols = ["total_bedrooms", "population",
    ↪ "households", "median_income", "rooms_per_household"], outputCol="features")
transformed = assembler.transform(df2)

```

```

[7]: from pyspark.ml.linalg import DenseVector
    #RDD
    dataRdd = transformed.select(transformed.median_house_value_final, transformed.
    ↪ features).rdd.map(tuple)
    data_df = dataRdd.toDF()
    data_df = data_df.withColumnRenamed("_1", "label").withColumnRenamed("_2",
    ↪ "features")
    data_df.show(5)

```

```

+-----+-----+
|label|      features|
+-----+-----+
|4.526|[129.0,322.0,126...|
|3.585|[1106.0,2401.0,11...|
|3.521|[190.0,496.0,177...|
|3.413|[235.0,558.0,219...|
|3.422|[280.0,565.0,259...|
+-----+-----+
only showing top 5 rows

```

We want to standardize the features, but not the response variable.

```

[8]: from pyspark.ml.feature import StandardScaler

```

```

[9]: # Feature scaling

# Initialize the `standardScaler`
standardScaler = StandardScaler(inputCol="features",
    ↪ outputCol="features_scaled",
                                withStd=True, withMean=False)

# Fit the DataFrame to the scaler; this computes the mean, standard deviation
    ↪ of each feature
scaler = standardScaler.fit(transformed)

# Transform the data in `df2` with the scaler
scaled_df = scaler.transform(transformed)

```

Split data into train set (80%), test set (20%) using seed=314

```
[10]: seed = 314
train_test = [0.8, 0.2]
train_data, test_data = scaled_df.randomSplit(train_test,seed)
```

Initialize the linear regression object with given parameters (1 POINT)

```
[11]: from pyspark.ml.regression import LinearRegression # note this is from the ML
      ↪ package

maxIter=10
regParam=0.3
elasticNetParam=0.8

lr =
      ↪ LinearRegression(featuresCol="features_scaled",labelCol="median_house_value_final",
      ↪ maxIter = 10, elasticNetParam = 0.8, regParam = 0.3)
```

Fit the model using the training data

```
[12]: model = lr.fit(train_data)
```

For each datapoint in the test set, make a prediction (hint: apply `transform()` to the model). You will want the returned object to be a dataframe

```
[13]: prediction = model.transform(test_data)
```

COMPUTE MSE (1 POINT)

Evaluate the model by computing Mean Squared Error (MSE), which is the average sum of squared differences between predicted and label.

This can be computed in a single line using `reduce()`

```
[14]: from pyspark.ml.evaluation import RegressionEvaluator
eval = RegressionEvaluator(labelCol = "median_house_value_final", predictionCol=
      ↪ "prediction",
                        metricName = "rmse")
mse = eval.evaluate(prediction, {eval.metricName: "mse"})
print("MSE: %.3f" % mse)
```

MSE: 0.755

RESULTS SECTION

```
[15]: # Code for Part 2

from pyspark.ml.feature import VectorAssembler
# subset the dataframe on these predictors
assembler2 = VectorAssembler(inputCols = [ "total_bedrooms", "population",
      ↪ "households", "median_income", "rooms_per_household", "housing_median_age"],
      ↪ outputCol="features")
```

```

transformed2 = assembler2.transform(df2)

standardScaler2 = StandardScaler(inputCol="features",
    ↳outputCol="features_scaled",
                                withStd=True, withMean=False)

# Fit the DataFrame to the scaler; this computes the mean, standard deviation
↳of each feature
scaler2 = standardScaler2.fit(transformed2)

# Transform the data in `df2` with the scaler
scaled_df2 = scaler2.transform(transformed2)

train_data2, test_data2 = scaled_df2.randomSplit(train_test,seed)
lr =
    ↳LinearRegression(featuresCol="features_scaled",labelCol="median_house_value_final",
    ↳maxIter = 10, elasticNetParam = 0.8, regParam = 0.3)
model2 = lr.fit(train_data2)
prediction2 = model2.transform(test_data2)
eval = RegressionEvaluator(labelCol = "median_house_value_final", predictionCol=
    ↳"prediction",
                                metricName = "rmse")
mse2 = eval.evaluate(prediction2, {eval.metricName: "mse"})
print("MSE: %.3f" % mse2)

```

MSE: 0.755

```

[16]: # Code for Part 3

df3 = df2.withColumn("households_per_pop", (df2.households/ df2.population))

assembler3 = VectorAssembler(inputCols = ["total_bedrooms", "population",
    ↳"households",
                                "median_income", "rooms_per_household", "housing_median_age",
    ↳"households_per_pop"], outputCol= "features")
transformed3 = assembler3.transform(df3)

standardScaler3 = StandardScaler(inputCol="features",
    ↳outputCol="features_scaled",
                                withStd=True, withMean=False)

scaler3 = standardScaler3.fit(transformed3)

scaled_df3 = scaler3.transform(transformed3)

train_data3, test_data3 = scaled_df3.randomSplit(train_test,seed)

```

```

lr3 = LinearRegression(featuresCol = 'features_scaled',
    ↳labelCol='median_house_value_final',maxIter=10, regParam=0.3,
    ↳elasticNetParam=0.8)
model3 = lr3.fit(train_data3)
prediction3 = model3.transform(test_data3)

# calculate MSE
mse3 = eval.evaluate(prediction3, {eval.metricName: 'mse'})

print("MSE: %.3f" % mse3)

```

MSE: 0.719

```

[17]: # Code for Part 4

lr4 =
    ↳LinearRegression(featuresCol="features_scaled",labelCol="median_house_value_final",
    ↳maxIter = 10, elasticNetParam = 1.0, regParam = 0.3)
model4 = lr4.fit(train_data)
prediction4 = model4.transform(test_data)
eval = RegressionEvaluator(labelCol = "median_house_value_final", predictionCol=
    ↳ "prediction",
                                metricName = "rmse")
mse4 = eval.evaluate(prediction4, {eval.metricName: "mse"})
print("MSE: %.3f" % mse4)

```

MSE: 0.775

Print dataframe containing question_part, MSE values for parts 1-4 in the next cell.

```

[28]: # print dataframe containing question_part, MSE
answer = spark.createDataFrame(
    [
        (1, mse),
        (2, mse2),
        (3, mse3),
        (4, mse4)
    ],
    ['question_part', 'MSE']
)
answer.show()

```

```

+-----+-----+
|question_part|          MSE|
+-----+-----+
|          1|0.7551749809615158|
|          2|0.7551747847051158|

```

```
|          3|0.7185818220171625|  
|          4|0.7746803974273968|  
+-----+-----+-----+
```

```
[30]: # Save notebook as PDF document  
!jupyter nbconvert --to pdf `pwd`/*.ipynb
```

```
[NbConvertApp] Converting notebook  
/home/jovyan/assignments/M5_4/cal_housing.ipynb to pdf  
[NbConvertApp] Writing 45772 bytes to ./notebook.tex  
[NbConvertApp] Building PDF  
[NbConvertApp] Running xelatex 3 times: ['xelatex', './notebook.tex', '-quiet']  
[NbConvertApp] Running bibtex 1 time: ['bibtex', './notebook']  
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no  
citations  
[NbConvertApp] PDF successfully created  
[NbConvertApp] Writing 56180 bytes to  
/home/jovyan/assignments/M5_4/cal_housing.pdf
```